

```
In [12]: import os
import random
from imutils import paths
import matplotlib.pyplot as plt
import argparse
import cv2
import tensorflow as tf
from tensorflow.keras.models import load_model
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
import numpy as np
from tensorflow import keras
```

```
In [13]: data = []
labels = []
imagePaths = list(paths.list_images("gesture"))
for imagePath in imagePaths:
    image = cv2.imread(imagePath)
    data.append(image)
    # extract the label from the image path and update the labels list
    label = imagePath.split(os.path.sep)[-1][0]

    labels.append(int(label))
```

```
In [3]: data = np.array(data, dtype="float") / 255.0
labels = np.array(labels)
```

```
In [4]: labels
```

```
Out[4]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
               1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2,  
               2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
               2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3, 3, 3, 3,  
               3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3,  
               3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,  
               4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,  
               4, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
               5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 6, 6, 6, 6,  
               6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6,  
               6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 7, 7, 7, 7, 7, 7, 7, 7,  
               7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7,  
               7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 7])
```

```
In [5]: (trainX, testX, trainY, testY) = train_test_split(data,
    labels, test size=0.15, random state=22)
```

```
In [48]: plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(trainX[i], cmap=plt.cm.gray)
    plt.xlabel(trainY[i])
plt.show()
```



```
In [7]: model = keras.Sequential([
    keras.layers.InputLayer(input_shape=(400, 400, 3)),
    keras.layers.Reshape(target_shape=(400, 400, 3)),
    keras.layers.Conv2D(filters=32, kernel_size=(3, 3), activation=tf.nn.relu),
    keras.layers.Conv2D(filters=64, kernel_size=(3, 3), activation=tf.nn.relu),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),
    keras.layers.Dropout(0.25),
    keras.layers.Flatten(),
    keras.layers.Dense(10)
])

# Define how to train the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

# Train the digit classification model
model.fit(trainX, trainY, epochs=5)
```

```
Epoch 1/5
8/8 [=====] - 140s 18s/step - loss: 31.4515 - accuracy: 0.1030
Epoch 2/5
8/8 [=====] - 121s 15s/step - loss: 2.8156 - accuracy: 0.3305
Epoch 3/5
8/8 [=====] - 117s 15s/step - loss: 0.9184 - accuracy: 0.6953
Epoch 4/5
8/8 [=====] - 117s 15s/step - loss: 0.4348 - accuracy: 0.8412
Epoch 5/5
8/8 [=====] - 117s 15s/step - loss: 0.1656 - accuracy: 0.9700
```

```
Out[7]: <tensorflow.python.keras.callbacks.History at 0x28568e57908>
```

In [8]: `model.summary()`

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
reshape_1 (Reshape)	(None, 400, 400, 3)	0
conv2d_2 (Conv2D)	(None, 398, 398, 32)	896
conv2d_3 (Conv2D)	(None, 396, 396, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 198, 198, 64)	0
dropout_1 (Dropout)	(None, 198, 198, 64)	0
flatten_1 (Flatten)	(None, 2509056)	0
dense_1 (Dense)	(None, 10)	25090570
=====		
Total params: 25,109,962		
Trainable params: 25,109,962		
Non-trainable params: 0		

In [10]: `test_loss, test_acc = model.evaluate(testX, testY)`

`print('Test accuracy:', test_acc)`

2/2 [=====] - 2s 815ms/step - loss: 0.3660 - accuracy: 0.9048

Test accuracy: 0.9047619104385376

```
In [15]: def get_label_color(val1, val2):
    if val1 == val2:
        return 'black'
    else:
        return 'red'

# Predict the labels of digit images in our test dataset.
predictions = model.predict(testX)

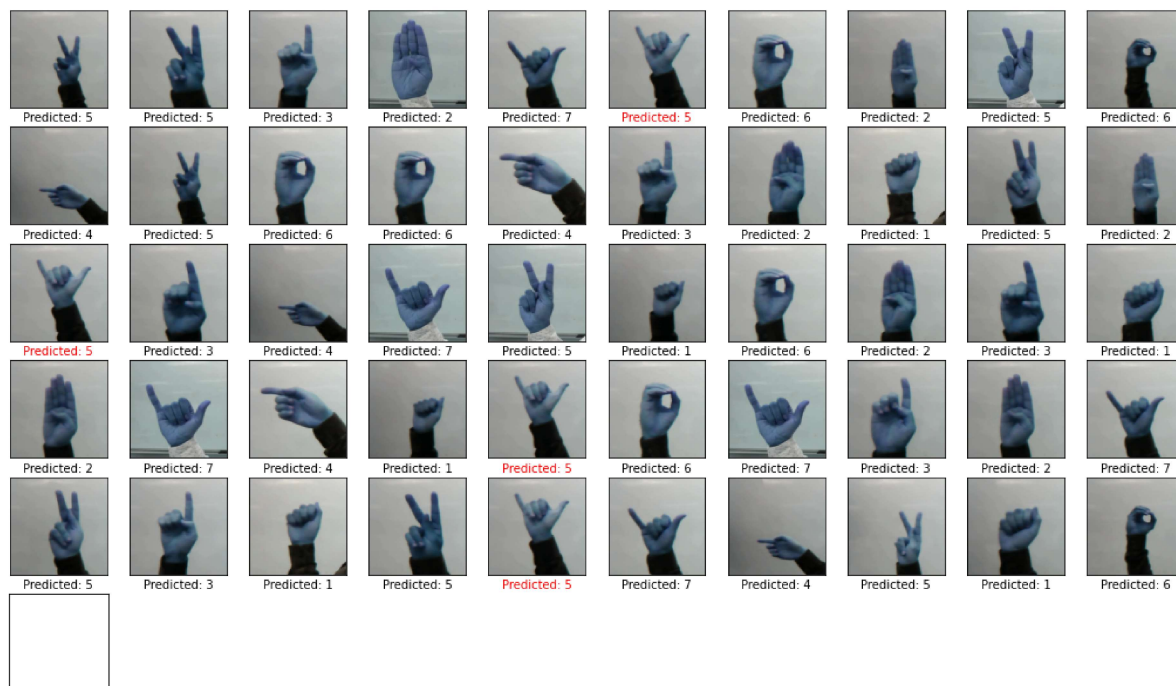
prediction_digits = np.argmax(predictions, axis=1)

# dataset, we will highlight it in red color.
plt.figure(figsize=(18, 18))
for i in range(100):
    ax = plt.subplot(10, 10, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    image_index = random.randint(0, len(prediction_digits))
    plt.imshow(testX[image_index], cmap=plt.cm.gray)
    ax.xaxis.label.set_color(get_label_color(prediction_digits[image_index],\
                                              testY[image_index]))
    plt.xlabel('Predicted: %d' % prediction_digits[image_index])
plt.show()
```

---

```
IndexError                                Traceback (most recent call last)
<ipython-input-15-ec4c7ded4c22> in <module>
    20 plt.grid(False)
    21 image_index = random.randint(0, len(prediction_digits))
--> 22 plt.imshow(testX[image_index], cmap=plt.cm.gray)
    23 ax.xaxis.label.set_color(get_label_color(prediction_digits[image_inde
x],\
    24                                     testY[image_index]))

IndexError: index 42 is out of bounds for axis 0 with size 42
```



```
In [17]: # Convert Keras model to TF Lite format.
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_float_model = converter.convert()
```

```
# Show model size in KBs.
float_model_size = len(tflite_float_model) / 1024
print('Float model size = %dKBs.' % float_model_size)
```

Float model size = 98087KBs.

```
In [18]: converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_quantized_model = converter.convert()

# Show model size in KBs.
quantized_model_size = len(tflite_quantized_model) / 1024
print('Quantized model size = %dKBs,' % quantized_model_size)
print('which is about %d%% of the float model size.' \
      % (quantized_model_size * 100 / float_model_size))
```

Quantized model size = 24526KBs,  
which is about 25% of the float model size.

```

In [27]: def evaluate_tflite_model(tflite_model):
    # Initialize TFLite interpreter using the model.
    interpreter = tf.lite.Interpreter(model_content=tflite_model)
    interpreter.allocate_tensors()
    input_tensor_index = interpreter.get_input_details()[0]["index"]
    output = interpreter.tensor(interpreter.get_output_details()[0]["index"])

    # Run predictions on every image in the "test" dataset.
    prediction_digits = []
    for test_image in trainX:
        # Pre-processing: add batch dimension and convert to float32 to match with
        # the model's input data format.
        test_image = np.expand_dims(test_image, axis=0).astype(np.float32)
        interpreter.set_tensor(input_tensor_index, test_image)

        # Run inference.
        interpreter.invoke()

        # Post-processing: remove batch dimension and find the digit with highest
        # probability.
        digit = np.argmax(output()[0])
        prediction_digits.append(digit)

    # Compare prediction results with ground truth labels to calculate accuracy.
    accurate_count = 0
    for index in range(len(prediction_digits)):
        if prediction_digits[index] == trainY[index]:
            accurate_count += 1
    accuracy = accurate_count * 1.0 / len(prediction_digits)

    return accuracy

# Evaluate the TF Lite float model. You'll find that its accuracy is identical
# to the original TF (Keras) model because they are essentially the same model
# stored in different format.
float_accuracy = evaluate_tflite_model(tflite_float_model)
print('Float model accuracy = %.4f' % float_accuracy)

# Evaluate the TF Lite quantized model.
# Don't be surprised if you see quantized model accuracy is higher than
# the original float model. It happens sometimes :)
quantized_accuracy = evaluate_tflite_model(tflite_quantized_model)
print('Quantized model accuracy = %.4f' % quantized_accuracy)
print('Accuracy drop = %.4f' % (float_accuracy - quantized_accuracy))

```

```

Float model accuracy = 0.9957
Quantized model accuracy = 0.9957
Accuracy drop = 0.0000

```

```

In [28]: f = open('gesture.tflite', "wb")
    f.write(tflite_quantized_model)
    f.close()

```

In [ ]: